n-Level Graph Partitioning*

Vitaly Osipov[†] Peter Sanders[‡]

Abstract

We present a multi-level graph partitioning algorithm based on the extreme idea to contract only a single edge on each level of the hierarchy. This obviates the need for a matching algorithm and promises very good partitioning quality since there are very few changes between two levels. Using an efficient data structure and new flexible ways to break local search improvements early, we obtain an algorithm that scales to large inputs and produces the best known partitioning results for many inputs. For example, in Walshaw's well known benchmark tables we achieve 155 improvements dominating the entries for large graphs.

1 Introduction

Many important applications of computer science involve processing large graphs, e.g., stemming from finite element methods, digital circuit design, route planning, social networks, etc. Very often these graphs need to be partitioned or clustered such that there are few edges between the blocks (pieces).

A successful heuristic for partitioning large graphs is the multilevel graph partitioning approach (MGP) depicted in Figure 1 where the graph is recursively contracted to a smaller graph with the same basic structure. After applying an initial partitioning algorithm to this small graph, the contraction is undone and, at each level, a local refinement method improves the partition induced by the coarser level. Section 2 explains the method in more detail. Most systems instantiate MGP in a very similar way: Maximal matchings are contracted between two levels that try to include as many heavy edges as possible. Local refinement uses a linear time variant of local search. MGP has two crucial advantages over most other approaches to graph partitioning: We get near linear execution time since the graph shrinks geometrically and we get good partitioning quality since a good solution on some level yields a good initial solution on the next finer level, i.e., local search needs little work to further improve the solution.

Our central idea is to get even better partitions by making subsequent levels as similar as possible – we (un)contract only a single edge between two levels. We call this n-GP

^{*}Partially supported by DFG grant SA 933/3-2.

[†]Universität Karlsruhe (TH), Germany, osipov@ira.uka.de

[‡]Universität Karlsruhe (TH), Germany, sanders@ira.uka.de

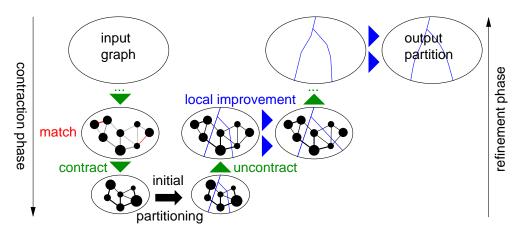


Figure 1: Multilevel graph partitioning.

since we have (almost) n levels of hierarchy. More details are described in Section 3. n-GP has the additional advantage that there is no longer a need for an algorithm finding heavy matchings. This is remarkable insofar as a considerable amount of work on approximate maximum weight matching was motivated by the MGP application [24, 5, 22, 18]. Still, at first glance, n-GP seems to have substantial disadvantages also. Firstly, storing each level explicitly would lead to quadratic space consumption. We avoid this by using a dynamic graph data structure with little space overhead. Secondly, choosing maximal matchings instead of just a single edge for contraction has the side effect that the graph is contracted everywhere, leading to a more uniform distribution of node weights. We solve this problem by explicitly factoring node weights into the edge rating function prioritizing the edges to be contracted. Already in [13, 14] edge ratings have proven to lead to better results for graph partitioning. Perhaps the most serious problem is that the most common approach to local search is to let it run for a number of steps proportional to the current number of nodes. In the context of n-GP this could lead to a quadratic overall number of local search steps. Therefore, we develop a new, more adaptive stopping criteria for the local search that drastically accelerates n-GP without significantly reducing partitioning quality.

We have implemented n-GP in the graph partitioner KaSPar (Karlsruhe Sequential Partitioner). Experiments reported in Section 5 indicate that KaSPar scales well to large networks, computes the best known partitions for many instances of a "standard benchmark" and needs time comparable to system that previously computed the best results for large networks. Section 6 summarizes the results and discusses future directions.

More Related Work

There has been a huge amount of research on graph partitioning so that we refer to introductory and overview papers such as [7, 8, 26, 30] for more material. Well-known software packages based on MGP are Chaco [12], DiBaP [19], Jostle [29, 30], Metis [16, 17], Party [23, 25], and Scotch [20, 21].

KaSPar was developed partly in parallel with KaPPa (Karlsruhe Parallel Partitioner)

[14]. KaPPa is a "classical" matching based MGP algorithm designed for scalable parallel execution and its local search only considers independent pairs of blocks at a time. Still, for k = 2, its interesting to compare KaSPar and KaPPa since KaPPa achieves the previously best partitioning results for many large graphs, since both systems use a similar edge ratings, and since running times for a two processor parallel code and a sequential code could be expected to be roughly comparable.

There is a long tradition of n-level algorithms in geometric data structures based on randomized incremental construction (e.g, [11, 1]). Our motivation for studying n-level are contraction hierarchies [10], a preprocessing technique for route planning that is at the same time simpler and an order of magnitude more efficient than previous techniques using a small number of levels.

2 Preliminaries

Consider an undirected graph $G=(V,E,c,\omega)$ with edge weights $\omega:E\to\mathbf{R}_{>0}$, node weights $c:V\to\mathbf{R}_{\geq 0},\ n=|V|,$ and m=|E|. We extend c and ω to sets, i.e., $c(V'):=\sum_{v\in V'}c(v)$ and $\omega(E'):=\sum_{e\in E'}\omega(e).\ \Gamma(v):=\{u:\{v,u\}\in E\}$ denotes the neighbors of v.

We are looking for blocks of nodes V_1, \ldots, V_k that partition V, i.e., $V_1 \cup \cdots \cup V_k = V$ and $V_i \cap V_j = \emptyset$ for $i \neq j$. The balancing constraint demands that $\forall i \in 1..k : c(V_i) \leq L_{\max} := (1 + \epsilon)c(V)/k + \max_{v \in V} c(v)$ for some parameter ϵ . The last term in this equation arises because each node is atomic and therefore a deviation of the heaviest node has to be allowed. The objective is to minimize the total $cut \sum_{i < j} w(E_{ij})$ where $E_{ij} := \{\{u, v\} \in E : u \in V_i, v \in V_j\}$. By default, our initial inputs will have unit edge and node weights. However, even those will be translated into weighted problems in the course of the algorithm.

Contracting an edge $\{u, v\}$ means replacing the nodes u and v by a new node x connected to the former neighbors of u and v. We set c(x) = c(u) + c(v). If replacing edges of the form $\{u, w\}, \{v, w\}$ would generate two parallel edges $\{x, w\}$, we insert a single edge with $\omega(\{x, w\}) = \omega(\{u, w\}) + \omega(\{v, w\})$. Uncontracting an edge e undoes its contraction. Partitions computed for the contracted graph are extrapolated to the uncontracted graph in the obvious way, i.e., u and v are put into the same block as x.

Local Search is done by moving single nodes between blocks. The gain $g_B(v)$ of moving node v to block B is decrease in total cut size caused by this move. For example, if v has 5 incident edges of unit weight, 2 of which are inside v's block and 3 of which lead to block b then $g_B(v) = 3 - 2 = 1$

3 *n*-Level Graph Partitioning

Figure 2 gives a high-level recursive summary of n-GP. The base case is some other partitioner used when the graph is sufficiently small. In KaSPar, contraction is stopped when either only 20k nodes remain, no further nodes are eligible for contraction, or there are less edges than nodes left. The latter happens when the graph consists of many independent components.

```
Function n\text{-}GP(G, k, \epsilon)

if G is small then return initialPartition(G, k, \epsilon)

pick the edge e = \{u, v\} with the highest rating

contract e; \mathcal{P} := n\text{-}GP(G, k, \epsilon); uncontract e

activate(u); activate(v); localSearch(v)

return \mathcal{P}
```

Figure 2: n-GP.

As observed in [14] Scotch [20] produces better initial partitions than metis, and therefore we also use it in KaSPar .

The edges to be contracted are chosen according to an edge rating function. KaSPar adopts the rating function

expansion*
$$(\{u, v\}) := \frac{\omega(\{u, v\})}{c(u)c(v)}$$

which fared best in [14]. As a further measure to avoid unbalanced inputs to the initial partitioner, KaSPar never allows a node v to participate in a contraction if the weight of v exceeds 1.5n/(20k). Selecting contracted edges can be implemented efficiently by keeping the contractable *nodes* in a priority queue sorted by the rating of their most highly rated incident edge.

In order to make contraction and uncontraction efficient, we use a "semidynamic" graph data structure: When contracting an edge $\{u,v\}$, we mark both u and v as deleted, introduce a new node w, and redirect the edges incident to u and v to w. The advantage of this implementation is that edges adjacent to a node are still stored in adjacency arrays which are more efficient than linked lists needed for a full fledged dynamic graph data structure. A disadvantages of our approach is a certain space overhead. However, it is relatively easy to show that this space overhead is bounded by a logarithmic factor even if we contract edges in some random fashion (see [4]). In Section 5 we will demonstrate experimentally that the overhead is actually often a small constant factor. Indeed, this is not very surprising since the edge rating function is not random, but designed to keep the contracted graph sparse. Overall, with respect to asymptotic memory overhead, n-GP is no worse than methods with a logarithmic number of levels.

3.1 Local Search Strategy

Our local search strategy is similar to the FM-algorithm [6] that is also used in many other MGP systems. We now outline our variant and then discuss differences.

Originally, all nodes are unmarked. Only unmarked nodes are allowed to be activated or moved from one block to another. Activating a node $v \in B'$ means that for blocks $\{B \neq B' : \exists \{v, u\} \in E \land u \in B\}$ we compute the gain

$$g_B(v) = \sum \{\omega(\{v, u\}) : \{v, u\} \in E, v \in B\} - \sum \{\omega(\{v, u\}) : \{v, u\} \in E, v \in B'\}$$

of moving v to block B. Node v is then inserted into the priority queue P_B using $g_B(v)$ as the priority. We call a queue P_B eligible if the highest gain node in P_b can be moved to block B without violating the balance constraint for block B. Local search repeatedly looks for the highest gain node v in any eligible priority queue P_B and moves v to block B. When this happens, node v becomes nonactive and marked, the unmarked neighbors of v get activated and the gains of the active neighbors are updated. The local search is stopped if either no eligible nonempty queues remain, or one of the stopping criteria described below applies. After the local search stops, it is rolled back to the lowest cut state reached during the search (which is the starting state if no improvement was achieved). Subsequently all previously marked nodes are unmarked. The local search is repeated until no improvement is achieved.

The main difference to the usual FM-algorithm is that our routine performs a highly localized search starting just at the uncontracted edge. Indeed, our local search does nothing if none of the uncontracted nodes is a border node, i.e., has a neighbor in another block. Other FM-algorithms initialize the search with all border nodes. In n-GP the local search may find an improvement quickly after moving a small number of nodes. However, in order to exploit this case, we need a way to stop the search much earlier than previous algorithms which limit the number of steps to a constant fraction of the current number of nodes |V|.

Stopping Using a Random Walk Model. It makes sense to make a stopping rule more adaptive by making it dependent on the past history of the search, e.g., on the difference between the current cut and the best cut achieved before.

We model the gain values in each step as identically distributed, independent random variables whose expectation μ and Variance σ^2 is obtained from the previously observed p steps. In Appendix A we show how from these (purely heuristical, i.e., technically unwarranted) assumptions we can derive that it is unlikely that the local search will produce a better cut if

$$p\mu^2 > \alpha\sigma^2 + \beta \tag{1}$$

where α and β are tuning parameters and μ is the average gain since the last improvement. For the variance σ^2 , we can also use the variance observed throughout the current local search. Parameter β is a base value that avoids stopping just after a small constant number of steps that happen to have small variance. Currently we set it to $\ln n$.

4 Trial Trees

It is a standard technique in optimization heuristics to improve results by repeating various parts of the algorithm. We generalize several approaches used in MGP by adapting an idea initially used in a fast randomized min-cut algorithm [15]: After reducing the number of nodes by a factor c, we perform two independent trials using different random seeds for tie breaking during contraction, initial partitioning, and local search. Among these trials the one with the smaller cut is used for continuing upwards. This way, we perform independent

trials at many levels of contraction controlled by a single tuning parameter c. As long as c > 2, the total number of contraction steps performed stays $\mathcal{O}(n)$.

5 Experiments

Implementation. We implemented KaSPar in C++ using gcc-4.3.2. We use priority queues based on paring heaps [28] available in the policy-based elementary data structures library (pb_ds) for implementing contraction and refinement procedures. In the following experimental study we compared KaSPar to Scotch 5.1, kMetis 4.0 and the same version of KaPPa as in [14].

System. We performed our experiments on a single core of an Intel Xeon Quad-core Processor featuring 2x4 MB of L2 cache and clocked at 2.667 GHz of a 2 processor Intel Xeon X5355 node with 16 GB of RAM running Suse Linux Enterprise 10.

Instances. We report results on two suites of instances summarized in Table 1. rggX is a random geometric graph with 2^X nodes that represent random points in the unit square and edges connect nodes whose Euclidean distance is below $0.55\sqrt{\ln n/n}$. This threshold was chosen in order to ensure that the graph is almost connected. DelaunayX is the Delaunay triangulation of 2^X random points in the unit square. Graphs bcsstk29..fetooth and ferotor..auto come from Chris Walshaw's benchmark archive [27]. Graphs bel, nld, deu and eur are undirected versions of the road networks of Belgium, the Netherlands, Germany, and Western Europe respectively, used in [3]. Instances af_shell9 and af_shell10 come from the Florida Sparse Matrix Collection [2]. coAuthorsDBLP, coPapersDBLP, citationCiteseer, coAuthorsCiteseer and cnr2000 are examples of social networks taken from [9].

For the number of partitions k we choose the values used in [27]: 2, 4, 8, 16, 32, 64. Our default value for the allowed imbalance is 3 % since this is one of the values used in [27] and the default value in Metis.

When not otherwise mentioned, we perform 10 repetitions for the small networks and 5 repetitions for the other. We report the arithmetic average of computed cut size, running time and the best cut found. When further averaging over multiple instances, we use the geometric mean in order to give every instance the same influence on the final figure.

Configuring the Algorithm. We use two sets of parameter settings fast and strong. These methods only differ in the constant factor α in the local search stopping rule, see Equation (1), in the contraction factor c for the trial tree (Section 4), and in the number of initial partitioning attempts a performed at the coarsest level of contraction:

\mathbf{S}	trategy	α	c	a
f	ast	1	8	$25/\log_2 k$
S	trong	4	2.5	$100/\log_2 k$

Table 1: Basic properties of the graphs from our benchmark set. The large instances are split into five groups: geometric graphs, FEM graphs, street networks, sparse matrices, and social networks. Within their groups, the graphs are sorted by size.

3.6.11	1	
Medium size	d instance	es
graph	n	m
rgg17	2^{17}	1457506
rgg18	2^{18}	3094566
Delaunay17	2^{17}	786352
Delaunay18	2^{18}	1572792
bcsstk29	13 992	605496
4elt	15606	91756
fesphere	16386	98 304
cti	16840	96464
memplus	17758	108384
cs4	33 499	87716
pwt	36519	289588
bcsstk32	44609	1970092
body	45087	327468
t60k	60005	178 880
wing	62032	243088
finan512	74752	522240
ferotor	99617	662431
bel	463514	1183764
nld	893 041	2279080
af_shell9	504855	17 084 020

Large instances		
graph	n	m
rgg20	2^{20}	13 783 240
Delaunay20	2^{20}	12582744
fetooth	78 136	905 182
598a	110971	1483868
ocean	143437	819 186
144	144649	2148786
wave	156317	2118662
m14b	214765	3358036
auto	448695	6629222
deu	4 378 446	10967174
eur	18029721	44435372
af_shell10	1508065	51 164 260
Social networks		
coAuthorCiteseer	227 320	1 628268
coAutorhDBLP	299067	1955352
cnr2000	325557	3216152
citationCiteseer	434102	32073440
coPaperDBLP	540486	30 491 458

Note that this are considerably less parameters compared to KaPPa. In particular, there is no need for selecting a matching algorithm, an edge coloring algorithm, or global and local iterations for refinement.

Scalability. Figure 3 shows the number of edges touched during contraction (KaSPar strong, small and large instances). We see that this scales linearly with the number of input edges and with a fairly small constant factor between 2 and 3. Interestingly, the number of local search steps during local improvement (Figure 4) decreases with increasing input size. This can be explained by the sublinear number of border vertices that we have in graphs that have small cuts and by small average search space sizes for the local search. Indeed, Figure 5 in the appendix indicates that the average length of local searches grows only logarithmically with n. All this translates into fairly complicated running time behavior. Still, Figure 6 in the appendix warrants the conclusion that running time scales "near linearly" with the input

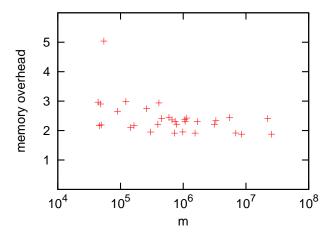


Figure 3: Number of edges created during contraction.

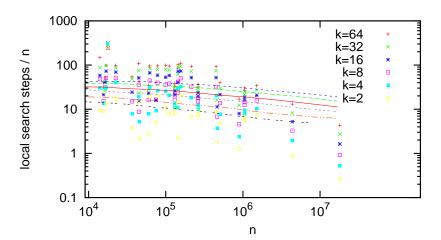


Figure 4: Total number of local search steps. The nearly straight lines represent series for the graphs rgg15..rgg24 and Delaunay15..Delaunay24 for different k.

size. The term in the running time depending on k grows sublinearly with the input size so that for very large graphs the number of blocks does not matter much.

Does the Random Walk Model Work? We have compared KaSPar fast with a variant where the stopping rule is disabled (i.e., $\alpha = \infty$). For the small instances this yields about 1 % better cut sizes at the cost of an order of magnitude larger running time. This is a small improvement both compared to the improvement KaSPar achieves over other systems and compared to just repeating KaSPar fast 10 times (see Table 2).

Do Trial trees help? We use the following evaluation: We run KaSPar strong and measure its elapsed time. Then for different values of initial partitionings a we repeat KaSPar strong without trial trees(c=0), until the sum of the run times of all repetitions exceeds the run time of KaSPar strong. Than for different values a we compare the best edge cut achieved during repeated runs to the one produced by KaSPar strong. Finally, we average the obtained results over 5 repetitions of this procedure. If we then quality the computed partitions, we usually get almost identical results (a fraction of a percent difference). However, most of the time trial trees are a bit better and for road networks we get considerable improvements. For example, for the European network we get an improvement of 10 % on average over all k.

Comparison with other Systems. Table 2 summarizes the results by computing geometric means over 10 runs for the small instances and over 5 runs for the large instances and social networks. We exclude the European road network for k=2 because KaPPa runs out of memory in that case. Detailed per instance results can be found in the appendix. KaPPa strong produces 5.9 % larger cuts than KasPar strong for small instances (average value) and 8.1 % larger cuts for the large instances. This comparison might seem a bit unfair because KaPPa is about five times faster. However, KaPPa is using k processors in parallel. Indeed, for k=2 KaSPar strong needs only about twice as much time. Also note that KaPPa strong needs about twice as much time as KaSPar fast while still producing 6 % larger cuts despite running in parallel. The case k=2 is also interesting because here KaPPa and KaSPar are most similar – parallelism does not play a big role (2 processors) and both local search strategies work only on two blocks at all time. Therefore 6 % improvement of KaSPar over KaPPa we can attribute mostly to the larger number of levels.

Scotch and kMetis are much faster than KaSPar but also produce considerably larger cuts – e.g., 32 % larger for large instances (kMetis, average). For the European road network, the difference in cut size even exceeds a factor of two. Such gaps usually cannot be breached by just running the faster solver a larger number of times. For example, for large instances, Scotch is only a factor around 4 faster than KaSPar fast, yet its best cut values obtained from 5 runs are still 12.7 % larger than the average values of KaSPar fast.

For social networks all systems have problems. KaSPar lags further behind in terms of speed but extends its lead with respect to the cut size. We mostly attribute the larger

¹This may not apply to the social networks which have considerably worse behavior.

Table 2: Geometric means over all instances.

code	sm	all grap	ohs	laı	rge graph	ns	soc	ial networ	rks
	best	avg.	t[s]	best	avg.	t[s]	best	avg.	t[s]
KaSPar strong	2675	2729	7.37	12450	12584	87.12	-	-	-
KaSPar fast	2717	2809	1.43	12655	12842	14.43	93657	99062	297.34
KaSPar fast, $\alpha = \infty$	2697	2780	23.21	-	-	-	-	-	-
KaPPa strong	2807	2890	2.10	13 323	13 600	28.16	117701	123613	78.00
KaPPa fast	2819	2910	1.29	13442	13727	16.67	117927	126914	46.40
kMetis	3097	3 348	0.07	15540	16656	0.71	117959	134803	1.42
Scotch	2926	3065	0.48	14475	15074	3.83	168764	168764	17.69

		La	arge Insta	ances					
k	Ka	SPar str	ong	Kal	aPPa strong				
	best	avg.	t[s]	best	avg.	t[s]			
2	2842	2873	36.89	2977	3054	15.03			
4	5642	5 707	60.66	6 190	6384	30.31			
8	10464	10580	75.92	11375	11652	37.86			
16	17345	17567	102.52	18678	19061	39.13			
32	27416	27 707	137.08	29156	29562	31.35			
64	41284	41570	170.54	43237	43644	22.36			

run time to the larger cut sizes relative to the number of nodes which greatly increase the number of local searches necessary. A further effect may be that the time for a local search step is proportional to the number of partitions adjacent to the nodes participating in the local search. For "well behaved" graphs this is mostly two, but for social networks which get denser on the coarser levels this value can get larger.

The Walshaw Benchmark [27] considers 34 graphs using $k \in \{2, 4, 8, 16, 32, 64\}$ and balance parameter $\epsilon \in \{0, 0.01, 0.03, 0.05\}$ giving a total of 816 table entries. Only cut sizes count – running time is not reported. We tried all combinations except the case $\epsilon = 0$ which KaSPar cannot handle yet. We ran KaSPar strong with a time limit of one hour and report the best result obtained in the appendix. KaSPar improved 155 values in the benchmark table: 42 for 1%, 49 for 3% and 64 for 5% allowed imbalance. Moreover, it reproduced equally sized cuts in 83 additional cases. If we count only results for graphs having over 44k nodes and $\epsilon > 0$, KaSPar improved 131 and reproduced 27 cuts, thus summing up to 63% of large graph table slots. We should note, that 51 of the new improvements are over partitioners different from KaPPa. Most of the improvements lie in the lower triangular part of the table, meaning that KaSPar is particularly good for either large graphs, or smaller graphs with small k. On the other hand, for small graphs, large k, and $\epsilon = 1\%$ KaSPar was often not able to obtain a feasible solution. A primary reason for this seems to be that

initial partitioning yields highly infeasible solutions that KaSPar is not able to to improve considerably during refinement. This is not astonishing, since Scotch targets $\epsilon = 3\%$ and does not even guarantee that.

6 Conclusion

n-GP is a graph partitioning approach that scales to large inputs and currently computes the best known partitions for many large graphs, at least when a certain imbalance is allowed. It is in some sense simpler than previous methods since no matching algorithm is needed. Although our current implementation of KaSPar is a considerable constant factor slower than the fastest available MGP partitioners, we see potential for further tuning. In particular, thanks to our adaptive stopping rule, KaSPar needs to do very little local search, in particular for large graphs and small k. Thus it suffices to tune the relatively simple contraction routine to obtain significant improvements. On the other hand, the adaptive stopping rule might also turn out to be useful for matching based MGP algorithms.

A lot of opportunities remain to further improve KaSPar. In particular, we did not yet attempt to handle the case $\epsilon = 0$ since this may require different local search strategies. We also want to try other initial partitioning algorithms and ways to integrate n-GP into other metaheuristics like evolutionary search.

We expect that n-GP could be generalized for other objective functions, for hypergraphs, and for graph clustering. More generally, the success of n-GP also suggests to look for more applications of the n-level paradigm.

An apparent drawback of n-GP is that it looks inherently sequential. However, we could probably obtain a good parallel algorithm by contracting small sets of highly rated, independent edges in parallel. Indeed, in the light of our results for KaSPar the complications coming from the need to find maximal matchings of heavy edges seem unnecessary, i.e., a parallelization of n-GP might be fast and simple.

Acknowledgements. We would like to thank Christian Schulz for supplying data for KaPPa, Scotch and Metis.

References

- [1] M. Birn, M. Holtgrewe, P. Sanders, and J. Singler. Simple and fast nearest neighbor search. In 11th Workshop on Algorithm Engineering and Experiments, 2010.
- [2] T. Davis. The University of Florida Sparse Matrix Collection, http://www.cise.ufl.edu/research/sparse/matrices, 2008.
- [3] D. Delling, P. Sanders, D. Schultes, and D. Wagner. Engineering route planning algorithms. In *Algorithmics of Large and Complex Networks*, volume 5515 of *LNCS State-of-the-Art Survey*, pages 117–139. Springer, 2009.

- [4] R. Dementiev, P. Sanders, D. Schultes, and J. Sibeyn. Engineering an external memory minimum spanning tree algorithm. In *IFIP TCS*, Toulouse, 2004.
- [5] D. Drake and S. Hougardy. Improved linear time approximation algorithms for weighted matchings. In 7th Workshop on Randomization and Approximation Techniques in Computer Science (APPROX), LNCS 2764, pages 14–23, 2003.
- [6] C. M. Fiduccia and R. M. Mattheyses. A Linear-Time Heuristic for Improving Network Partitions. In 19th Conf. on Design Automation, pages 175–181, 1982.
- [7] P. Fjallstrom. Algorithms for graph partitioning: A survey. Linkoping Electronic Articles in Computer and Information Science, 3(10), 1998.
- [8] V. K. G. Karypis. A fast and high quality multilevel scheme for partitioning irregular graphs. SIAM JOURNAL ON SCIENTIFIC COMPUTING, 20(1):359–392, 1998.
- [9] R. Geisberger, P. Sanders, and D. Schultes. Better approximation of betweenness centrality. In 10th Workshop on Algorithm Engineering and Experimentation, pages 90–108, San Francisco, 2008. SIAM.
- [10] R. Geisberger, P. Sanders, D. Schultes, and D. Delling. Contraction hierarchies: Faster and simpler hierarchical routing in road networks. In 7th Workshop on Experimental Algorithms, volume 5038 of LNCS, pages 319–333. Springer, 2008.
- [11] L. J. Guibas, D. E. Knuth, and M. Sharir. Randomized incremental construction of Delaunay and Voronoi diagrams. *Algorithmica*, 7(4):381–413, 1992.
- [12] B. Hendrickson. Chaco: Software for partitioning graphs. http://www.sandia.gov/~bahendr/chaco.html.
- [13] M. Holtgrewe. A scalable coarsening phase for a multi-level partitioning algorithm. Diploma thesis, Universität Karlsruhe, 2009.
- [14] M. Holtgrewe, P. Sanders, and C. Schulz. Engineering a scalable high quality graph partitioner. In 24th IEEE International Parallel and Distributed Processing Symposium, 2010. to appear, see also arXiv:0910.2004.
- [15] D. R. Karger and C. Stein. A new approach to the minimum cut problem. J. ACM, 43(4):601-640, 1996.
- [16] G. Karypis and V. Kumar. MeTis: Unstructured Graph Partitioning and Sparse Matrix Ordering System, Version 4.0. http://www.cs.umn.edu/~metis.
- [17] G. Karypis and V. Kumar. MeTiS, A Software Package for Partitioning Unstructured Graphs, Partitioning Meshes, and Computing Fill-Reducing Orderings of Sparse Matrices, Version 4.0, 1998.

- [18] J. Maue and P. Sanders. Engineering algorithms for approximate weighted matching. In 6th Workshop on Exp. Algorithms (WEA), volume 4525 of LNCS, pages 242–255. Springer, 2007.
- [19] H. Meyerhenke, B. Monien, and T. Sauerwald. A new diffusion-based multilevel algorithm for computing graph partitions. *Journal of Parallel and Distributed Computing*, 69(9):750–761, 2009.
- [20] F. Pellegrini. SCOTCH: Static Mapping, Graph, Mesh and Hypergraph Partitioning, and Parallel and Sequential Sparse Matrix Ordering Package, 2007. http://www.labri.fr/perso/pelegrin/scotch/.
- [21] F. Pellegrini. SCOTCH 5.1 User's guide. Technical report, Laboratoire Bordelais de Recherche en Informatique, Bordeaux, France, 2008.
- [22] S. Pettie and P. Sanders. A simpler linear time $2/3 \epsilon$ approximation for maximum weight matching. *Information Processing Letters*, 91(6):271–276, 2004.
- [23] R. Preis. PARTY Partitioning Library , 1996. http://www2.cs.uni-paderborn.de/cs/robsy/party.html.
- [24] R. Preis. Linear time 1/2-approximation algorithm for maximum weighted matching in general graphs. In 16th Symp. on Theoretical Aspects of Computer Science, volume 1563 of LNCS, pages 259–269. Springer, 1999.
- [25] R. Preis and R. Diekmann. The PARTY Partitioning Library, User Guide. Technical report, University of Paderborn, Germany, 1996. Tr-rsfb-96-02.
- [26] K. Schloegel, G. Karypis, and V. Kumar. Graph partitioning for high performance scientific simulations. Technical Report 00-018, University of Minnesota, 2000.
- [27] A. J. Soperm, C. Walshaw, and M. Cross. A combined evolutionary search and multilevel optimisation approach to graph partitioning. *J. Global Optimization*, 29(2):225-241, 2004. http://staffweb.cms.gre.ac.uk/~c.walshaw/partition/.
- [28] A. Tavory, V. Dreizin, and B. Kosnik. Policy-based data structures. http://gcc.gnu.org/onlinedocs/libstdc++/ext/pb_ds/ IBM Haifa and Redhat, 2004.
- [29] C. Walshaw. JOSTLE -graph partitioning software, 2005. http://staffweb.cms.gre.ac.uk/~wc06/jostle/.
- [30] C. Walshaw and M. Cross. JOSTLE: Parallel Multilevel Graph-Partitioning Software An Overview. In F. Magoules, editor, *Mesh Partitioning Techniques and Domain Decomposition Techniques*, pages 27–58. Civil-Comp Ltd., 2007. (Invited chapter).

A Derivation of Stopping Rules

Consider a situation where p steps of local search have been performed with average value μ and variance σ^2 . Then in the next s steps, we can expect a deviation from the expectation $(p+s)\mu$ by something of the order $\sqrt{s\sigma^2}$. The expression $(p+s)\mu + \sqrt{s\sigma^2}$ is maximized for $s^* := \frac{\sigma^2}{4\mu^2}$. Now the idea is to stop when for some tuning parameter x, $(p+s^*)\mu + x\sqrt{s^*\sigma^2} > 0$, i.e., it is reasonably likely that a random walk modelling our local search can still give an improvement. This translates to the condition $p > \frac{\sigma^2}{\mu^2}(\frac{x}{2} - \frac{1}{4})$ or simply $p\mu^2 \gg \sigma^2$.

B Additional Figures and Tables

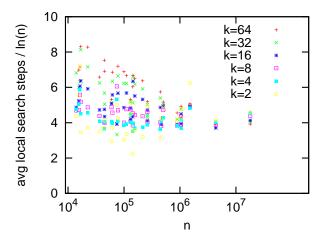


Figure 5: Average length of local searches.

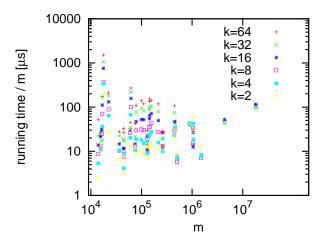


Figure 6: Running Time.

Graph	k		KaSPar fast	t	K	aPPa strong	S]	KaPPa fast		Kal	Pa minima	.l		scotch			metis	
	ĺ	best	avg	time	best	avg	time	best	avg	time	best	avg	time	best	avg	time	best	avg	time
coAuthorsCiteseer	2	17855	18003	25.77	21775	26462	12.93	29894	30997	11.98	32678	35492	6.75	34065	34065	5.34	21587	22674	0.30
coAuthorsCiteseer	4	34180	35315	51.75	43778	46540	28.43	44837	47156	17.03	50845	55514	5.78	52277	52277	7.51	39649	41560	0.33
coAuthorsCiteseer	8	49574	50054	85.49	56574	57647	46.35	53838	55686	25.77	61397	62752	5.10	69988	69988	9.61	56289	56996	0.36
coAuthorsCiteseer	16	59574	59915	124.51	66173	66648	55.26	62126	63085	31.04	65681	67007	5.71	83457	83457	11.41	68295	68744	0.39
coAuthorsCiteseer	32	67953	68752	169.78	72331	72736	64.53	71603	72062	30.34	74119	74760	5.49	90807	90807	12.92	77399	78254	0.41
coAuthorsCiteseer	64	76210	77326	193.85	77603	78756	64.45	79411	79872	26.47	81773	82244	5.82	100737	100737	14.20	84538	85426	0.44
citationCiteseer	2	32181	32247	49.44	35122	36696	24.37	34858	48466	15.31	47641	61055	11.51	37175	37175	5.87	33684	34344	0.67
citationCiteseer	4	67194	68371	135.82	76897	79782	47.87	76994	101369	27.40	120656	133916	12.54	79543	79543	11.02	73536	77524	0.76
citationCiteseer	8	103743	105663	297.70	119852	126129	85.92	118505	133337	47.09	188731	196204	13.41	124441	124441	15.37	108655	116082	0.83
citationCiteseer	16	148932	151256	507.69	156984	164984	114.26	156132	160555	57.02	218710	224851	14.18	163941	163941	18.83	153846	157000	0.91
citationCiteseer	32	198757	203173	841.73	205922	207923	147.02	198771	207089	111.12	248894	259090	45.56	210957	210957	21.88	197146	200650	0.98
citationCiteseer	64	255722	258037	1213.93	247462	248994	148.58	240660	241980	115.18	270692	279247	39.04	265971	265971	25.08	242010	244427	1.10
coAuthorsDBLP	2	45292	45650	82.42	54803	56140	23.13	55263	61619	16.96	63305	64872	11.78	63368	63368	8.23	48952	50341	0.53
coAuthorsDBLP	4	80408	81575	144.65	94651	97597	54.59	97007	98865	32.09	123373	126675	9.52	109856	109856	11.73	88513	88734	0.61
coAuthors <u>DB</u> LP	8	109940	113575	263.08	126261	128129	82.57	116839	118190	46.32	144839	147038	8.62	142749	142749	14.49	115201	117074	0.69
coAuthorsDBLP	16	132067	135259	440.42	144229	145229	98.64	137946	138968	46.24	152803	154368	7.51	169706	169706	16.97	138399	140149	0.75
coAuthorsDBLP	32	152146	154501	787.16	157754	159086	113.74	151883	153606	62.94	160331	161779	14.80	189201	189201	18.95	160842	161565	0.82
coAuthorsDBLP	64	168939	169122	1099.75	169681	170403	123.69	169283	169671	46.08	174708	175679	10.11	207486	207486	20.70	175660	177172	0.88
cnr2000	2	210	236	49.54	2597	3789	25.90	2430	4835	23.21	2422	5053	18.48	20537	20537	3.70	1773	2451	7.44
cnr2000	4	1569	1973	64.02	6089	6971	46.71	6284	6939	27.50	6175	7138	13.34	26809	26809	6.20	4301	6326	8.05
cnr2000	8	4096	4974	75.55	7914	8510	55.07	7378	7911	33.20	7684	8308	12.67	31373	31373	8.18	7899	16951	8.85
cnr2000	16	6943	14824	92.57	8784	10382	61.45	9399	9567	32.72	9805	10003	12.45	34967	34967	10.84	12601	81752	9.47
cnr2000	32	384058	400272	188.69	360661	363687	88.28	368182	372503	48.97	374786	375787	14.89	432813	432813	12.59	368062	409130	9.78
cnr2000	64	713772	723710	483.68	694270	700504	103.42	706366	712434	52.27	722754	727917	15.49	727685	727685	14.06	723221	737874	10.31
coPapersDBLP	2	462530	466947	372.39	512389	527205	80.25	490054	552438	66.76	528953	585647	60.58	622378	622378	42.06	599794	634286	2.33
coPapersDBLP	4	822518	838005	705.59	937267	952505	122.23	1021741	1034181	88.66	1409276	1428094	42.62	1188052	1188052	76.19	1073007	1091355	2.58
coPapersDBLP	8	1188694	1213398	1794.77	1257622	1293223	201.41	1296044	1315313	131.60	1690906	1751688	32.16	1685436	1685436	98.29	1442079	1495943	2.81
coPapersDBLP	16	1534078	1544591	3993.84	1540054	1571957	318.56	1593642	1614871	165.98	1816467	1852634	28.73	2028374	2028374	131.89	1864836	1886340	3.01
coPapersDBLP	32	1789129	1798109	6550.18	1828015	1850535	411.34	1790694	1861113	276.51	1926975	2009450	37.35	2380424	2380424	156.07	2087868	2122569	3.17
coPapersDBLP	64	2039271	2054249	10897.41	2164396	2177596	423.03	2051766	2061702	244.46	2132793	2139541	31.87	2697328	2697328	148.72	2341150	2347850	3.39

Table 3: All results for social network instances

Graph	k	K	aSPar stro	ng	I	KaSPar fas	t	K	aPPa stroi	ıg		KaPPa fast		Ka	PPa minin	nal		scotch			metis	
		best	avg	time	best	avg	time	best	avg	time	best	avg	time	best	avg	time	best	avg	time	best	avg	time
fe_tooth fe_tooth	2	3844 6937	3987 6999	5.86 8.54	3840 7034	3981 7146	1.16 1.49	3951 7012	4336 7189	3.75 5.22	3854 7126	4353 7757	2.44	4109 7780	6490 9157	1.59 0.96	4259 8304	4259 8304	0.38	4372 7805	4529 8280	0.08
fe_tooth	8	11482	11564	13.43	11574	12007	1.49	12272	12721	6.83	12215	12678	4.06	13243	13671	0.96	12999	12999	1.09	13334	13768	0.08
fe_tooth	16	17744	17966	21.24	17968	18288	2.79	18302	18570	7.18	18198	18524	3.55	19559	19813	0.65	20816	20816	1.59	20035	20386	0.09
fe_tooth	32	25888	26248	35.12	26249	26592	4.03	26397	26617	5.28	26404	26677	2.92	28070	28391	0.53	28430	28430	2.13	28547	29052	0.10
fe_tooth 598a	64	36259 2371	36469 2384	49.65 6.50	36741 2378	40385 2389	5.80 1.84	36862 2387	37002 2393	4.71 5.64	36795 2391	36992 2401	2.57 3.79	38423 2456	39095 2485	0.62 2.99	38401 2417	38401 2417	2.69 0.39	39233 2444	39381 2513	0.12
598a	4	7897	7921	11.15	7935	7977	2.42	8235	8291	10.24	8291	8385	5.94	9224	9862	2.62	8246	8246	0.95	8466	8729	0.14
598a	8	15929	15984	22.31	15992	16125	3.48	16502	16641	12.21	16461	16598	7.07	17351	17899	2.98	17490	17490	1.63	17170	17533	0.16
598a	16	26046	26270	38.39	26102	26672	5.05	26467	26825	17.74	26670	26887	12.51	27983	28596	6.76	29804	29804	2.37	27857	28854	0.17
598a 598a	32 64	39625 58362	40019 58945	60.60 87.52	40563 58326	40986 59199	7.25 10.72	40946 59148	41190 59387	18.16 14.15	40928 59026	41186 59233	11.91 9.64	43111 61396	43741 61924	7.74 6.21	44756 64561	44756 64561	3.21 4.11	43256 61888	44213 62703	0.19
fe_ocean	2	317	317	5.55	317	322	1.66	314	317	3.21	314	318	2.11	343	355	1.71	402	402	0.18	540	579	0.11
fe_ocean	4	1801	1810	9.40	1817	1837	1.95	1756	1822	6.30	1754	1822	3.03	1990	2051	1.10	2000	2000	0.44	2102	2140	0.11
fe_ocean fe_ocean	8 16	4044 7992	4097 8145	14.33 22.41	4084 8120	4195 8359	2.51 3.39	4104 8188	4252 8350	6.33 5.62	4143 8294	4330 8469	2.93 3.04	4689 9457	4987 9553	0.73	4956 9351	4956 9351	0.81 1.27	5256 10115	5472 10377	0.12
fe_ocean	32	13320	13518	36.53	13526	13806	5.00	13593	13815	4.34	13618	14042	2.15	15465	15657	0.47	15089	15089	1.83	16565	16877	0.15
fe_ocean	64	21326	21739	62.46	22059	22209	7.78	21636	21859	3.68	21809	21973	2.02	24147	24275	0.51	23246	23246	2.49	24198	24531	0.17
144	2	6455	6507	12.81	6461	6491	3.04	6559	6623	7.45	6563	6638	5.23	6747	6799	3.64	6695	6695	0.66	6804	6972	0.20
144 144	8	15312 25130	15471 25409	24.73 38.13	15717 25557	15774 26039	4.10 5.54	16870 26300	16963 26457	13.33 20.11	16998 26435	17122 26614	7.00 10.49	17364 27206	18101 27829	2.97 2.93	16899 28172	16899 28172	2.24	17144 28006	17487 28194	0.21
144	16	37872	38404	69.35	38830	39161	8.30	39010	39319	26.04	39266	39492	17.53	40264	41977	6.63	43712	43712	3.12	42861	43041	0.24
144	32	57082	57492	106.40	57353	57860	11.73	58331	58631	24.60	58175	58652	16.03	61774	62171	8.79	63224	63224	4.14	61716	62481	0.26
144	64	80313 8661	80770 8720	144.77 16.19	80609 8650	81293 8690	16.05 3.25	82286 8832	82452 9132	19.11 8.24	82029 8809	82493 9108	12.05 4.72	86067 8966	86950 9324	8.16 2.84	88246 9337	88246 9337	5.25 0.83	86534 9169	87208 9345	0.30
wave	4	16806	16920	29.56	16871	16978	4.39	17008	17250	14.51	17263	17503	6.84	18041	21189	1.81	19995	19995	1.72	19929	21906	0.19
wave	8	28681	28817	46.61	28865	29200	6.01	30690	31419	20.63	30628	31371	9.79	32617	33937	1.50	33357	33357	2.61	33223	33639	0.21
wave	16	42918	43208	75.97	43267	43770	8.31	44831	45048	20.54	44936	45202	10.73	46293	47270	1.47	48903	48903	3.53	48404	49000	0.22
wave wave	32 64	63025 87243	63159 87554	112.19 150.37	62764 87403	63266 87889	11.88 16.88	63981 88376	64390 88964	14.94 12.51	64004 88924	64532 89297	8.19 6.09	68085 92366	68620 93424	1.02	70581 96759	70581 96759	4.68 5.90	68062 92148	68604 94083	0.25
m14b	2	3828	3846	20.03	3845	3870	4.49	3862	3954	11.16	3900	3945	7.80	3951	4208	5.63	3872	3872	0.70	4036	4155	0.31
m14b	4	13015	13079	26.51	13111	13160	5.42	13543	13810	18.77	14104	14211	10.21	14990	15094	4.42	13484	13484	1.71	13932	14560	0.33
m14b	8 16	25573	25756	45.33	25839	26086	7.27 10.07	27330 45352	27393	24.97	27411	27450	11.76	28241	28517	3.78	27839	27839	2.86	28138 48314	28507 49269	0.34
m14b m14b	32	42212 66314	42458 66991	83.25 133.88	42727 66942	43365 68017	14.52	68107	45762 69075	28.11 29.94	45931 68715	46108 69223	19.27 17.99	48769 72484	49397 73598	6.41 8.12	50778 75453	50778 75453	4.25 5.72	72746	74135	0.40
m14b	64	99207	100014	198.23	99964	100666	20.91	101053	101455	25.26	101410	101861	17.46	106361	107173	10.24	109404	109404	7.38	107384	108141	0.44
auto	2	9740	9768	68.39	9744	9776	10.99	9910	10045	30.09	9863	10856	18.86	10313	11813	12.12	10666	10666	1.61	10781	12147	0.83
auto	8	25988 45099	26062 45232	75.60 97.60	26072 45416	26116 45806	13.35 15.98	28218 46272	29481 46652	64.01 85.89	29690 47163	29995 48229	33.11 46.36	32473 49447	33371 53617	8.93 8.42	29046 49999	29046 49999	3.52 5.42	27469 49691	30318 52422	0.86
auto	16	76287	76715	153.46	77376	77801	20.81	78713	79769	87.41	79711	80683	58.20	84236	86001	12.25	84462	84462	7.84	85562	89139	0.91
auto	32	121269	121862	246.50	122406	123052	28.12	124606	125500	71.77	124920	125876	46.44	131545	133723	20.23	133403	133403	10.58	133026	134086	0.99
doloupov p20	64	174612 1711	174914 1731	352.09 196.33	174712 1726	176214 1753	38.76 12.88	177038 1858	177595 1882	62.64 35.43	177461 1879	178119 1898	44.14 18.66	185836 1911	187424	25.39 13.87	193170 1874	193170	13.68	188555 2054	189699 2194	1.08
delaunay_n20 delaunay_n20	4	3418	3439	130.67	3460	3480	13.21	3674	3780	64.08	3784	3826	24.34	3857	1937 3900	8.97	3723	1874 3723	2.35	4046	4094	1.11
delaunay_n20	8	6278	6317	104.37	6364	6387	13.71	6670	6854	70.07	6688	6872	41.92	7161	7303	6.15	7180	7180	3.58	7705	8029	1.13
delaunay_n20	16	10183	10218	84.33	10230	10327	13.80	10816	11008	67.92	10882	11061	48.05	11307	11533	6.31	11266	11266	4.77	11854	12440	1.14
delaunay_n20 delaunay_n20	32 64	15905 23935	16026 23962	101.69 97.09	16211 24193	16236 24263	14.90 16.40	16813 24799	17086 25179	42.67 22.04	16814 24946	17150 25129	24.44 12.83	17993 26314	18179 27001	3.33 1.79	17784 26163	17784 26163	6.04 7.34	18816 28318	19304 28543	1.18
rgg_n_2_20_s0	2	2162	2201	198.61	2146	2217	16.75	2377	2498	33.24	2378	2497	24.66	2400	2530	20.71	2832	2832	1.41	3023	3326	1.57
rgg_n_2_20_s0	4	4323	4389	130.00	4382	4448	17.18	4867	5058	38.50	4870	4973	21.06	5114	5200	11.11	5737	5737	2.82	5786	6174	1.56
rgg_n_2_20_s0 rgg_n_2_20_s0	8 16	7745 12596	7915 12792	103.66 86.19	8031 12981	8174 13148	17.81 17.93	8995 14953	9391 15199	46.06 35.86	9248 15013	9493 15339	25.50 24.61	9426 15039	9632 15442	7.83 7.20	11251 17157	11251 17157	6.13	11365 17498	11771 18125	1.54
rgg_n_2_20_s0	32	20403	20478	100.03	20805	20958	18.99	23430	23917	26.04	23383	24222	16.93	23842	24164	3.94	28078	28078	7.96	27765	28495	1.58
rgg_n_2_20_s0	64	30860	31066	97.83	31203	31584	20.50	34778	35354	11.62	35086	35539	9.95	35252	35629	2.09	38815	38815	9.83	41066	42465	1.58
af_shell10	2	26225	26225	317.11	26225	26225	37.00	26225	26225	78.65	26225	26225	65.31	26525	26640	59.76	26825	26825	3.64	27625	28955	2.99
af_shell10 af_shell10	8	55075 97709	55345 100233	210.61 179.51	55875 100325	56375 102667	36.59 38.47	54950 101425	55265 102335	91.96 136.99	54950 102125	55500 103180	51.52 61.16	58366 110369	58627 111081	22.11 16.03	58500 105375	58500 105375	7.60 11.97	61100 117650	64705 120120	3.04
af_shell10	16	163125	165770	212.12	163600	165360	40.47	165025	166427	106.63	165625	166480	69.97	174677	175918	17.00	171725	171725	16.45	184350	188765	3.04
af_shell10	32	248268	252939	191.53	252555	256262	43.14	253525	255535	80.85	252487	255746	52.00	270249	275149	9.25	269375	269375	21.66	289400	291590	3.13
af_shell10 deu	64	372823 167	376512	207.76 231.47	378031 175	382191 179	49.38 58.31	379125	382923 221	43.01	380225 230	384140 240	29.43 47.55	400378 233	404085	4.82 38.11	402275 295	402275 295	27.33 3.19	421285	427047 286	3.18 5.38
deu	2	419	172 426	244.12	427	447	58.84	214 533	542	68.20 76.87	531	545	47.55	544	243 580	25.65	726	726	6.46	268 699	761	5.35
deu	8	762	773	250.50	781	792	59.20	922	962	99.76	935	973	45.05	974	1007	19.57	1235	1235	9.84	1174	1330	5.24
deu	16	1308	1333	278.31	1332	1387	61.82	1550	1616	105.96	1556	1618	78.82	1593	1656	21.79	2066	2066	13.11	2041	2161	5.19
deu deu	32 64	2182 3610	2217 3631	283.79 293.53	2251 3679	2295 3737	62.50 64.38	2548 4021	2615 4093	73.17 49.55	2535 4078	2641 4146	41.93 31.63	2626 4193	2711 4317	11.50 5.97	3250 4978	3250 4978	16.28 19.41	3319 5147	3445 5385	5.28
eur	2	133	138	1946.34	162	211	792.68	4021	4033	43.00	4010	4140	01.00	4130	4011	5.51	469	469	12.45	0141	5565	0.01
eur	4	355	375	2168.10	416	431	794.41	543	619	441.11	580	646	223.96	657	697	113.35	952	952	25.37	846	1626	29.40
eur	8	774	786	2232.31	823	834	809.21	986	1034	418.29	1013	1034	207.41	1060	1119	80.92	1667	1667	38.67	1675	3227	29.04
eur	16 32	1401 2595	1440 2643	2553.40 2598.84	1575 2681	1597 2761	930.59 958.24	1760 3186	1900 3291	497.93 417.52	1907 3231	1935 3314	295.81 306.52	1931 3202	2048 3386	94.56 55.63	2922 4336	2922 4336	51.50 65.16	3519 7424	9395 9442	30.58 30.81
eur	64	4502	4526	2533.56	4622	4675	868.75	5290	5393	308.17	5448	5538	183.98	5569	5770	29.64	6772	6772	77.14	11313	12738	30.30

Table 4: All results for large instances.

Graph	2	?	4	:	8	3	10	6	35	2	6	4
add20	641	594	1212	1177	1814	1704	2427	2121		2687		3236
data	190	188	405	383	699	660		1162		1865		2885
3elt	90	89	201	199	361	342	654	569		969		1564
uk	19	19	41	42	92	84	179	152		258		438
add32	10	10	33	33	66	66	117	117	212	212		493
bcsstk33	10105	10097	21756	21508	34377	34178	56687	54860		78132		108505
whitaker3	126	126	382	380	670	656	1163	1093		1717		2567
crack	184	183	370	362	696	678	1183	1092		1707		2566
wing_nodal	1703	1696	3609	3572	5574	5443	8624	8422		11980		16134
fe_4elt2	130	130	349	349	622	605	1051	1014		1657		2537
vibrobox	11538	10310	19267	19199	25190	24553	35514	32167	46331	41399		49521
bcsstk29	2818	2818	8035	8159	14212	13965	23808	21768		34886		57054
4elt	138	138	325	321	561	534	1009	939		1559		2596
fe_sphere	386	386	798	768	1236	1152	1914	1730		2565		3663
cti	318	318	950	944	1815	1802	3056	2906	5044	4223		5875
memplus	5698	5489	10234	9559	12599	11785	14410	13241	16340	14395		16857
cs4	378	367	970	940	1520	1467	2285	2206	3521	3090		4169
bcsstk30	6347	6335	16617	16622	34761	34604	72028	71234		115770		173945
bcsstk31	2723	2701	7351	7444	13371	13417	24791	24277	42745	38086		60528
fe_pwt	340	340	704	704	1441	1442	2835	2806		5612		8454
bcsstk32	4667	4667	9247	9492	20855	21490	37372	37673	72471	61144		95199
fe_body	262	262	599	636	1079	1156	1858	1931		3202		5282
t60k	78	75	213	211	470	465	866	849	1493	1391		2211
wing	803	787	1683	1666	2616	2589	4147	4131	6271	5902		8132
brack2	708	708	3027	3038	7144	7269	11969	11983	18496	17798		26557
finan512	162	162	324	324	648	648	1296	1296	2592	2592		10560
fe_tooth	3819	3823	6938	7103	11650	11935	18115	18283	26604	25977		35980
fe_rotor	2055	2045	7405	7480	12959	13165	21093	20773	33588	32783		47461
598a	2390	2388	7992	8154	16179	16467	26196	26427	40513	40674		59098
fe_ocean	388	387	1856	1878	4251	4299	8276	8432	13841	13660		21548
144	6489	6479	15196	15345	25455	25818	38940	39352	58359	58126		81145
wave	8716	8682	16891	17475	29207	30511	43697	44611	64198	64551		88863
m14b	3828	3826	13034	13391	25921	26666	42513	43975	67990	67770		101551
auto	10004	10042	26941	27790	45731	47650	77618	79847	123296	124991	179309	175975

Table 5: Walshaw Benchmark with $\epsilon=1$

Graph	2		4		8	3	1	6	32	2	64	1
add20	636	576	1195	1158	1765	1690	2331	2095	2862	2493		3152
data	186	185	379	378	662	650	1163	1133	1972	1802		2809
3elt	87	87	199	198	346	336	587	565	1035	958	1756	1542
uk	18	18	40	40	84	81	158	148	281	251	493	414
add32	10	10	33	33	66	66	117	117	212	212	509	493
bcsstk33	10064	10064	21083	21035	34150	34078	55372	54510	80548	77672	113269	107012
whitaker3	126	126	381	378	662	655	1125	1092	1757	1686	2733	2535
crack	182	182	360	360	685	676	1132	1082	1765	1679	2739	2553
wing_nodal	1681	1680	3572	3561	5424	5401	8476	8316	12282	11938	16891	15971
fe_4elt2	130	130	349	343	607	598	1022	1007	1686	1633	2658	2527
vibrobox	11538	10310	19239	18778	24691	24171	34226	31516	43532	39592	52242	49123
bcsstk29	2818	2818	7983	8045	14041	13817	22448	21410	35660	34407	58644	55366
4elt	137	137	319	319	533	523	942	914	1631	1537	2728	2581
fe_sphere	384	384	792	764	1193	1152	1816	1706	2715	2477	3965	3547
cti	318	318	924	917	1724	1716	2900	2778	4396	4132	6330	5763
memplus	5626	5355	10145	9418	12521	11628	14168	13130	15850	14264	18364	16724
cs4	366	361	959	936	1490	1467	2215	2126	3152	3048	4479	4169
bcsstk30	6251	6251	16497	16537	34275	34513	70851	70278	117500	114005	178977	171727
bcsstk31	2676	2676	7183	7181	13090	13246	24211	23504	39298	37459	60847	58667
fe_pwt	340	340	704	704	1416	1419	2787	2784	5649	5606	8557	8346
bcsstk32	4667	4667	8778	8799	20035	21023	35788	36613	61485	59824	96086	92690
fe_body	262	262	598	601	1033	1054	1767	1800	2906	2947	4982	5212
t60k	71	71	211	207	461	454	851	822	1423	1391	2264	2198
wing	789	774	1660	1636	2567	2551	4034	4015	6005	5832	8316	8043
brack2	684	684	2853	2839	6980	6994	11622	11741	17491	17649	26679	26366
finan512	162	162	324	324	648	648	1296	1296	2592	2592	10635	10560
fe_tooth	3794	3792	6862	6946	11422	11662	17655	17760	25685	25624	35962	35830
fe_rotor	1960	1963	7182	7222	12546	12852	20356	20521	32114	31763	47613	47049
598a	2369	2367	7873	7955	15820	16031	25927	25966	39525	39829	58101	58454
fe_ocean	311	311	1710	1698	3976	3974	7919	7838	12942	12746	21217	21033
144	6456	6438	15122	15250	25301	25491	37899	38478	56463	57354	80621	80767
wave	8640	8616	16822	16936	28664	28839	42620	43063	62281	62743	86663	87325
m14b	3828	3823	12977	13136	25550	26057	42061	42783	65879	67326	98188	100286
auto	9716	9782	25979	26379	45109	45525	76016	77611	120534	122902	172357	174904

Table 6: Walshaw Benchmark with $\epsilon=3$

Graph	,	2	4	:	8	,	1	6	35	2	64	4
add20	610	550	1186	1157	1755	1675	2267	2081	2786	2463	3270	3152
data	183	181	369	368	640	628	1130	1086	1907	1777	3073	2798
3elt	87	87	198	197	336	330	572	560	1009	950	1645	1539
uk	18	18	39	40	81	78	150	139	272	246	456	410
add32	10	10	33	33	63	65	117	117	212	212	491	493
bcsstk33	9914	9914	20198	20584	33971	33938	55273	54323	79159	77163	111659	106886
whitaker3	126	126	380	378	658	650	1110	1084	1741	1686	2663	2535
crack	182	182	361	360	673	667	1096	1080	1749	1679	2681	2548
wing_nodal	1672	1668	3541	3536	5375	5350	8419	8316	12149	11879	16566	15873
fe_4elt2	130	130	340	335	596	583	1013	991	1665	1633	2608	2516
vibrobox	11538	10310	19021	18778	24203	23930	34298	31235	42890	39592	50994	48200
bcsstk29	2818	2818	7936	7942	13619	13614	21914	20924	34906	33818	57220	54935
4elt	137	137	318	315	519	516	925	902	1574	1532	2673	2565
fe_sphere	384	384	784	764	1219	1152	1801	1692	2678	2477	3904	3547
cti	318	318	900	890	1708	1716	2830	2725	4227	4037	6127	5684
memplus	5516	5267	10011	9299	12458	11555	14047	13078	15749	14170	18213	16454
cs4	363	356	955	936	1483	1467	2184	2126	3115	2995	4394	4116
bcsstk30	6251	6251	16186	16332	34146	34350	69520	70043	114960	113321	175723	170591
bcsstk31	2676	2676	7099	7152	12941	13058	23603	23254	38150	37459	60768	57534
fe_pwt	340	340	700	701	1405	1409	2772	2777	5545	5546	8410	8310
bcsstk32	4622	4644	8454	8481	19678	20099	35208	35965	60441	59824	94238	91006
fe_body	262	262	596	601	1017	1054	1723	1784	2807	2887	4834	4888
t60k	65	65	202	196	457	454	839	818	1398	1376	2229	2168
wing	784	770	1654	1636	2528	2551	3998	4015	5915	5806	8228	7991
brack2	660	660	2745	2739	6671	6781	11358	11558	17256	17529	26321	26281
finan512	162	162	324	324	648	648	1296	1296	2592	2592	10583	10560
fe_tooth	3780	3773	6825	6864	11337	11662	17404	17603	25216	25624	35466	35476
fe_rotor	1950	1955	7052	7045	12380	12566	20039	20132	31450	31576	46749	46608
598a	2338	2336	7763	7851	15544	15721	25585	25808	39144	39369	57412	58031
fe_ocean	311	311	1705	1697	3946	3941	7618	7722	12720	12746	20886	20667
144	6373	6362	15036	15250	25025	25259	37433	38225	56345	56926	79296	80257
wave	8598	8563	16662	16820	28615	28700	42482	42800	61788	62520	85658	86663
m14b	3806	3802	12976	13136	25292	25679	41750	42608	65231	66793	98005	99063
auto	9487	9450	25399	25883	44520	45039	75066	76488	120001	122378	171459	173968

Table 7: Walshaw Benchmark with $\epsilon=5$